

# Eedris\_Busari\_Fraugster\_Assignment

November 24, 2020

## 1 QUESTION ONE (1)

## 2 DATA LOADING

```
[1]: # I would have required to use dask which is a parallel form of data loading if
# the size of the data were heavier to increase time efficiency and avoiding
↳ loading
# all the data into the memory. An alternative is to chunk the data but it is
↳ not as efficient, comparatively
# because of the concatenation required at the end of the chunk process.
import pandas as pd
data=pd.read_csv('realestate_fraugster_case.csv',sep=';',index_col=False)
data.head(10)
```

```
[1]:
```

	street	city	zip	state	beds	baths	\
0	3526 HIGH ST	SACRAMENTO	95838	CA	2.0	1.0	
1	51 OMAHA CT	sacramento	95823	CA	3.0	1.0	
2	2796 BRANCH ST	SACRAMENTO	95815	CA	2.0	1.0	
3	2805 JANETTE WAY	SACRAMENTO	95815	CA	2.0	1.0	
4	6001 MCMAHON DR	SACRAMENTO	95824	CA	2.0	1.0	
5	5828 PEPPERMILL CT	SACRAMENTO	95841	CA	3.0	1.0	
6	6048 OGDEN NASH WAY	SACRAMENTO	95842	CA	3.0	2.0	
7	2561 19TH AVE	SACRAMENTO	95820	CA	3.0	1.0	
8	11150 TRINITY RIVER DR Unit 114	RANCHO CORDOVA	95670	CA	2.0	2.0	
9	7325 10TH ST	RIO LINDA	95673	CA	3.0	2.0	

  

	sq_ft	type	sale_date	price	latitude	longitude
0	836.0	Residential	1943-01-09 11:56:01	59222	38.631913	-121.434879
1	1167.0	Residential	1996-11-08 23:09:38	68212	38.478902	-121.431028
2	796.0	Residential	1915-01-05 07:31:45	68880	38.618305	-121.443839
3	852.0	Residential	1998-10-22 04:46:05	69307	38.616835	-121.439146
4	797.0	Residential	1972-01-05 20:52:32	81900	38.51947	-121.435768
5	1122.0	Condo	1918-01-13 23:10:18	89921	38.662595	-121.327813
6	1104.0	Residential	1949-06-16 12:35:50	90895	38.681659	-121.351705
7	1177.0	Residential	1971-01-31 00:55:56	91002	38.535092	-121.481367
8	941.0	Condo	1955-12-30 14:44:20	94905	38.621188	-121.270555
9	1146.0	Residential	1977-06-03 09:55:18	98937	38.700909	-121.442979

### 3 QUESTION TWO (2)

#### 4 DATA CLEANING

The data cleaning steps would be done in three phases as:

#### 5 PHASE 1: THE GENERAL OUTLOOK AND PROFILE OF THE DATASET

##### 6 (a) Statistical Description

The “describe” method of panda’s dataframe gives the statistical description of the dataset. This helps to see the count of unique values, most frequent value, how the values deviate or vary from one another percentile, among others.

```
[2]: data.describe(include='all')
```

```
[2]:
```

	street	city	zip	state	beds	baths	\
count	985	985	985	985	985.000000	985.000000	
unique	981	42	69	3	NaN	NaN	
top	7 CRYSTALWOOD CIR	SACRAMENTO	95648	CA	NaN	NaN	
freq	2	437	72	983	NaN	NaN	
mean	NaN	NaN	NaN	NaN	2.911675	1.776650	
std	NaN	NaN	NaN	NaN	1.307932	0.895371	
min	NaN	NaN	NaN	NaN	0.000000	0.000000	
25%	NaN	NaN	NaN	NaN	2.000000	1.000000	
50%	NaN	NaN	NaN	NaN	3.000000	2.000000	
75%	NaN	NaN	NaN	NaN	4.000000	2.000000	
max	NaN	NaN	NaN	NaN	8.000000	5.000000	

	sq__ft	type	sale_date	price	latitude	\
count	985.000000	985	986	985	985	
unique	NaN	7	986	606	969	
top	NaN	Residential	1999-05-22	19:42:16	4897	38.423251
freq	NaN	914	1	49	5	
mean	1314.916751	NaN	NaN	NaN	NaN	
std	853.048243	NaN	NaN	NaN	NaN	
min	0.000000	NaN	NaN	NaN	NaN	
25%	952.000000	NaN	NaN	NaN	NaN	
50%	1304.000000	NaN	NaN	NaN	NaN	
75%	1718.000000	NaN	NaN	NaN	NaN	
max	5822.000000	NaN	NaN	NaN	NaN	

longitude

```

count          985
unique          967
top    -121.444489
freq           5
mean           NaN
std            NaN
min            NaN
25%            NaN
50%            NaN
75%            NaN
max            NaN

```

```
[3]: data.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 986 entries, 0 to 985
Data columns (total 12 columns):
#   Column      Non-Null Count  Dtype
---  -
0   street      985 non-null    object
1   city        985 non-null    object
2   zip         985 non-null    object
3   state       985 non-null    object
4   beds        985 non-null    float64
5   baths       985 non-null    float64
6   sq__ft      985 non-null    float64
7   type        985 non-null    object
8   sale_date   986 non-null    object
9   price       985 non-null    object
10  latitude    985 non-null    object
11  longitude   985 non-null    object
dtypes: float64(3), object(9)
memory usage: 92.6+ KB

```

## 7 (b) Data Type Formats

When trying to convert to specific datatypes, the rows that do not comply to the rules of this datatype are identified as errors. These would help in making suitable corrections on the identified observations.

Also, possible operations on the columns depend on the datatype. The correct datatypes would also help to identify errors in the columns. In this section, emphasis would be made on the numeric columns while the non-numeric features would form the basis for the Inconsistency check in phase two

The above information could help determine the need for type conversion. The columns with 'object' datatypes need to be investigated to determine which ones would require conversion

## 8 i) The 'city', 'state', 'street' and 'type' object columns are non-numeric values

The 'city', 'state', and 'type' look tempting to convert to the category dtypes for memory efficiency and optimization. However, they would be left as object because the dataset is not large enough to cause memory issues. Also, if converted to category dtype, the addition of new distinct value into the columns would generate 'NaN' error.

## 9 ii) The 'sale\_date' column being a date would be converted to date datatype.

```
data['sale_date'] = pd.to_datetime(data.sale_date, format='%Y-%m-%d %H:%M:%S')
```

```
data['sale_date'] = pd.to_datetime(data.sale_date, format='%Y-%m-%d %H:%M:%S')
```

Running the above line gives errors such as the one identified below

ValueError: time data 1917-07-24 08:12:24% doesn't match format specified

```
[4]: # The error causing rows were identified and corrected as follows
data["sale_date"].replace({"2013-12-19 04:05:22A": "2013-12-19 04:05:22",
    ↳ "1917-07-24 08:12:24%": "1917-07-24 08:12:24", "1918-02-25 20:36:13%":
    ↳ "1918-02-25 20:36:13"}, inplace=True)
```

## 10 iii) The 'zip' and 'price' object columns have numeric values. These are supposed to be integer values. This is checked and the rows with errors are identified

```
[5]: # The error causing rows were identified in the zip column
for j, value in enumerate(data['zip']):
    try:
        int(value)
    except ValueError:
        print('The identified error index {}: {}'.format(j, value))
```

The identified error index 30: '957f58'

The identified error index 985: nan

```
[6]: # The error causing rows were identified in the price column
for j, value in enumerate(data['price']):
    try:
        int(value)
    except ValueError:
        print('The identified error index {}: {}'.format(j, value))
```

The identified error index 115: '298000D'

The identified error index 985: nan

```
[7]: # The typographical error were corrected intuitively as follows
data["zip"].replace({"957f58": "95758"}, inplace=True)
data["price"].replace({"298000D": "298000"}, inplace=True)
```

11 iv) The 'longitude' and 'latitude' object columns have floating values. These are checked and the rows with errors identified

```
[8]: for j, value in enumerate(data['longitude']):
    try:
        float(value)
    except ValueError:
        print('Index error for Longitude {}: {}'.format(j, value))
```

```
Index error for Longitude 121: '-121.2286RT'
Index error for Longitude 147: '-121.363757$'
```

```
[9]: for j, value in enumerate(data['latitude']):
    try:
        float(value)
    except ValueError:
        print('Index error for Latitude {}: {}'.format(j, value))
```

```
Index error for Latitude 109: '38.410992C'
```

```
[10]: # The typographical error were replaced intuitively as follows
data["longitude"].replace({"-121.2286RT": "-121.228678", "-121.363757$": "-121.
↪363757"}, inplace=True)
data["latitude"].replace({"38.410992C": "38.410992"}, inplace=True)
```

```
[11]: #data = data.astype({'longitude': 'float64', 'latitude': 'float64', 'price':
↪'int64', 'zip': 'int64'})
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 986 entries, 0 to 985
Data columns (total 12 columns):
#   Column      Non-Null Count  Dtype
---  -
0   street      985 non-null    object
1   city        985 non-null    object
2   zip         985 non-null    object
3   state       985 non-null    object
4   beds        985 non-null    float64
5   baths       985 non-null    float64
6   sq__ft      985 non-null    float64
7   type        985 non-null    object
```

```

8   sale_date  986 non-null    object
9   price      985 non-null    object
10  latitude   985 non-null    object
11  longitude  985 non-null    object
dtypes: float64(3), object(9)
memory usage: 92.6+ KB

```

## 12 PHASE 2: THE INCONSISTENCY CHECK

(The data consistency check is used for the following:

Redundancy such as duplicates, irrelevant datapoints, format error among others in both the columns and the rows

To do this, we check the consistency of non-numeric features (type, state, city and street) by:

(i) Capitalization Consistency Check

```

[12]: #The solution to the inconsistency in the case format (lower and upper cases)
      ↳ can be solved by either making all the letters
      # The upper case would be used in this case
data= data.apply(lambda x: x.astype(str).str.upper() if x.name in ['street',
      ↳ 'type', 'city', 'state'] else x)
data

```

```

[12]:
      street      city  zip state  beds  baths  sq_ft  \
0      3526 HIGH ST  SACRAMENTO  95838    CA    2.0    1.0    836.0
1       51 OMAHA CT  SACRAMENTO  95823    CA    3.0    1.0   1167.0
2     2796 BRANCH ST  SACRAMENTO  95815    CA    2.0    1.0    796.0
3    2805 JANETTE WAY  SACRAMENTO  95815    CA    2.0    1.0    852.0
4    6001 MCMAHON DR  SACRAMENTO  95824    CA    2.0    1.0    797.0
..      ...      ...      ...      ...      ...      ...
981    6932 RUSKUT WAY  SACRAMENTO  95823    CA    3.0    2.0   1477.0
982   7933 DAFFODIL WAY  CITRUS HEIGHTS  95610    CA    3.0    2.0   1216.0
983    8304 RED FOX WAY    ELK GROVE  95758    CA    4.0    2.0   1685.0
984  3882 YELLOWSTONE LN  EL DORADO HILLS  95762    CA    3.0    2.0   1362.0
985          NAN          NAN      NaN    NaN    NaN    NaN    NaN

```

```

      type      sale_date  price  latitude  longitude
0  RESIDENTIAL  1943-01-09 11:56:01  59222  38.631913  -121.434879
1  RESIDENTIAL  1996-11-08 23:09:38  68212  38.478902  -121.431028
2  RESIDENTIAL  1915-01-05 07:31:45  68880  38.618305  -121.443839
3  RESIDENTIAL  1998-10-22 04:46:05  69307  38.616835  -121.439146
4  RESIDENTIAL  1972-01-05 20:52:32  81900  38.51947  -121.435768
..      ...      ...      ...      ...      ...
981  RESIDENTIAL  1955-09-26 15:13:23  234000  38.499893  -121.45889
982  RESIDENTIAL  1950-09-13 20:59:20  235000  38.708824  -121.256803
983  RESIDENTIAL  1933-04-10 20:13:38  235301  38.417  -121.397424

```

```

984 RESIDENTIAL 1934-06-05 04:16:37 235738 38.655245 -121.075915
985          NAN 1940-02-11 05:29:17      NaN      NaN      NaN

```

[986 rows x 12 columns]

## (ii) Duplicate Row Check

```

[13]: # Duplicate row check would result into repetition with no new information in
      ↪ the dataset.
      # Therefore, observations that have been earlier recorded should be deleted. It
      ↪ could happen as a result of double submission
      # file merging, among others
      data.drop_duplicates(inplace=True)
      data

```

```

[13]:
      street      city  zip state  beds  baths  sq_ft \
0      3526 HIGH ST  SACRAMENTO  95838    CA    2.0    1.0   836.0
1       51 OMAHA CT  SACRAMENTO  95823    CA    3.0    1.0  1167.0
2     2796 BRANCH ST  SACRAMENTO  95815    CA    2.0    1.0   796.0
3    2805 JANETTE WAY  SACRAMENTO  95815    CA    2.0    1.0   852.0
4    6001 MCMAHON DR  SACRAMENTO  95824    CA    2.0    1.0   797.0
..      ...      ...      ...      ...      ...      ...
981    6932 RUSKUT WAY  SACRAMENTO  95823    CA    3.0    2.0  1477.0
982   7933 DAFFODIL WAY  CITRUS HEIGHTS  95610    CA    3.0    2.0  1216.0
983   8304 RED FOX WAY    ELK GROVE  95758    CA    4.0    2.0  1685.0
984  3882 YELLOWSTONE LN  EL DORADO HILLS  95762    CA    3.0    2.0  1362.0
985          NAN      NAN      NaN    NAN    NaN    NaN    NaN

```

```

      type      sale_date  price  latitude  longitude
0  RESIDENTIAL 1943-01-09 11:56:01  59222  38.631913 -121.434879
1  RESIDENTIAL 1996-11-08 23:09:38  68212  38.478902 -121.431028
2  RESIDENTIAL 1915-01-05 07:31:45  68880  38.618305 -121.443839
3  RESIDENTIAL 1998-10-22 04:46:05  69307  38.616835 -121.439146
4  RESIDENTIAL 1972-01-05 20:52:32  81900  38.51947 -121.435768
..      ...      ...      ...      ...      ...
981 RESIDENTIAL 1955-09-26 15:13:23  234000  38.499893 -121.45889
982 RESIDENTIAL 1950-09-13 20:59:20  235000  38.708824 -121.256803
983 RESIDENTIAL 1933-04-10 20:13:38  235301    38.417 -121.397424
984 RESIDENTIAL 1934-06-05 04:16:37  235738  38.655245 -121.075915
985          NAN 1940-02-11 05:29:17      NaN      NaN      NaN

```

[986 rows x 12 columns]

## (iii) Irrelevant/Redundant Row Check

```

[14]: # Since, it is a real estate sales data. Some columns could be seen as unique
      ↪ identifiers.

```

```

# Unavailability or missingness of this identifiers would render the
↳ observation(row) redundant
# An identifier here would be the Longitude and Latitude.
# This is because the house/bed/baths sold would not be identified without this
↳ information.
# Therefore, rows with this missing values should be removed
import numpy as np
data = data.dropna(axis=0, subset=['longitude','latitude'])
data

```

```

[14]:
      street      city  zip state  beds  baths  sq_ft  \
0      3526 HIGH ST    SACRAMENTO  95838    CA    2.0    1.0   836.0
1       51 OMAHA CT    SACRAMENTO  95823    CA    3.0    1.0  1167.0
2     2796 BRANCH ST    SACRAMENTO  95815    CA    2.0    1.0   796.0
3    2805 JANETTE WAY    SACRAMENTO  95815    CA    2.0    1.0   852.0
4    6001 MCMAHON DR    SACRAMENTO  95824    CA    2.0    1.0   797.0
..      ...      ...      ...      ...      ...      ...
980   9169 GARLINGTON CT    SACRAMENTO  95829    CA    4.0    3.0  2280.0
981    6932 RUSKUT WAY    SACRAMENTO  95823    CA    3.0    2.0  1477.0
982   7933 DAFFODIL WAY  CITRUS HEIGHTS  95610    CA    3.0    2.0  1216.0
983    8304 RED FOX WAY      ELK GROVE  95758    CA    4.0    2.0  1685.0
984  3882 YELLOWSTONE LN  EL DORADO HILLS  95762    CA    3.0    2.0  1362.0

```

```

      type      sale_date  price  latitude  longitude
0  RESIDENTIAL  1943-01-09 11:56:01  59222  38.631913  -121.434879
1  RESIDENTIAL  1996-11-08 23:09:38  68212  38.478902  -121.431028
2  RESIDENTIAL  1915-01-05 07:31:45  68880  38.618305  -121.443839
3  RESIDENTIAL  1998-10-22 04:46:05  69307  38.616835  -121.439146
4  RESIDENTIAL  1972-01-05 20:52:32  81900   38.51947  -121.435768
..      ...      ...      ...      ...      ...
980 RESIDENTIAL  1951-06-03 12:20:20  232425  38.457679  -121.35962
981 RESIDENTIAL  1955-09-26 15:13:23  234000  38.499893  -121.45889
982 RESIDENTIAL  1950-09-13 20:59:20  235000  38.708824  -121.256803
983 RESIDENTIAL  1933-04-10 20:13:38  235301   38.417  -121.397424
984 RESIDENTIAL  1934-06-05 04:16:37  235738  38.655245  -121.075915

```

[985 rows x 12 columns]

#### (iv) Typographical and Format Errors

The unique values of the non-numeric columns ('type', 'state', 'city', and 'street') as shown in Out[2]: above are free text, which is prone to typographical error and human discretion in its format used. A look at the unique values show these errors.

As can be seen in the state column, there are typographical error as 'CA', 'CA3', 'CA-' is pointing to a singular state 'CA'.

[1] The solution to the 'states' column can be either of:



- a) Delete the column since it is a single-valued column and would not help in any ML modelling task.
- b) Correct the spelling and typo-errors.

For completeness of the dataset, I will just replace the values with 'CA'

```
[15]: # Check the unique values in the 'state' column and also save a copy of the
      ↪ data with a new name
      print(data.state.unique())
      new_data=data.copy()
```

```
['CA' 'CA3' 'CA-']
```

```
[16]: #new_data.loc[new_data['state'] == 'CA']
      new_data=data.loc[data['state'] == 'CA']
      new_data.state.unique()
```

```
[16]: array(['CA'], dtype=object)
```

[2] The solution to the 'type' column:

```
[31]: #The unique values in the type column are replaced appropriately
      new_data.type.unique()
      new_data["type"].replace({"RESIDENTIAL%": "RESIDENTIAL","RESIDEN_TIAL":
      ↪ "RESIDENTIAL","RESIDENTIAL)": "RESIDENTIAL"}, inplace=True)
      new_data.type.unique()
```

```
[31]: array(['RESIDENTIAL', 'CONDO', 'MULTI-FAMILY', 'UNKOWN'], dtype=object)
```

[3] The solution to the 'city' column:

```
[18]: # To check the count and unique values in the column
      print(new_data.city.nunique())
      new_data.city.unique()
```

41

```
[18]: array(['SACRAMENTO', 'RANCHO CORDOVA', 'RIO LINDA', 'CITRUS HEIGHTS',
            'NORTH HIGHLANDS', 'ANTELOPE', 'SACRAMENTO@', 'ELK GROVE',
            'ELVERTA', 'GALT', 'CARMICHAEL', 'ORANGEVALE', 'FOLSOM',
            'ELK GROVE<>', 'MATHER', 'POLLOCK PINES', 'GOLD RIVER',
            'EL DORADO HILLS', 'RANCHO MURIETA', 'WILTON', 'GREENWOOD',
            'FAIR OAKS', 'CAMERON PARK', 'LINCOLN', 'PLACERVILLE',
            'MEADOW VISTA', 'ROSEVILLE', 'ROCKLIN', 'AUBURN', 'LOOMIS',
            'EL DORADO', 'PENRYN', 'GRANITE BAY', 'FORESTHILL',
            'DIAMOND SPRINGS', 'SHINGLE SPRINGS', 'COOL', 'WALNUT GROVE',
            'GARDEN VALLEY', 'SLOUGHHOUSE', 'WEST SACRAMENTO'], dtype=object)
```

```
[19]: # One way to do this is to create a list of valid cities in California
# Then, check the "city" column with this list.
# Any value that is present in the 'city' column but not available in the actual
# city list would be investigated
actual_city=['SACRAMENTO', 'RANCHO CORDOVA', 'RIO LINDA', 'CITRUS HEIGHTS',
            'NORTH HIGHLANDS', 'ANTELOPE', 'ELK GROVE',
            'ELVERTA', 'GALT', 'CARMICHAEL', 'ORANGEVALE', 'FOLSOM',
            'ELK GROVE', 'MATHER', 'POLLOCK PINES', 'GOLD RIVER',
            'EL DORADO HILLS', 'RANCHO MURIETA', 'WILTON', 'GREENWOOD',
            'FAIR OAKS', 'CAMERON PARK', 'LINCOLN', 'PLACERVILLE',
            'MEADOW VISTA', 'ROSEVILLE', 'ROCKLIN', 'AUBURN', 'LOOMIS',
            'EL DORADO', 'PENRYN', 'GRANITE BAY', 'FORESTHILL',
            'DIAMOND SPRINGS', 'SHINGLE SPRINGS', 'COOL', 'WALNUT GROVE',
            'GARDEN VALLEY', 'SLOUGHHOUSE', 'WEST SACRAMENTO']
check_this= new_data[~new_data.city.isin(actual_city)].city
check_this
```

```
[19]: 28    SACRAMENTO@
      76    ELK GROVE<>
      Name: city, dtype: object
```

```
[20]: #The unique values in the type column are replaced appropriately
new_data["city"].replace({"SACRAMENTO@": "SACRAMENTO", "ELK GROVE<>": "ELK_
↳ GROVE"}, inplace=True)
print(new_data.city.unique())
new_data.city.unique()
```

39

```
[20]: array(['SACRAMENTO', 'RANCHO CORDOVA', 'RIO LINDA', 'CITRUS HEIGHTS',
            'NORTH HIGHLANDS', 'ANTELOPE', 'ELK GROVE', 'ELVERTA', 'GALT',
            'CARMICHAEL', 'ORANGEVALE', 'FOLSOM', 'MATHER', 'POLLOCK PINES',
            'GOLD RIVER', 'EL DORADO HILLS', 'RANCHO MURIETA', 'WILTON',
            'GREENWOOD', 'FAIR OAKS', 'CAMERON PARK', 'LINCOLN', 'PLACERVILLE',
            'MEADOW VISTA', 'ROSEVILLE', 'ROCKLIN', 'AUBURN', 'LOOMIS',
            'EL DORADO', 'PENRYN', 'GRANITE BAY', 'FORESTHILL',
            'DIAMOND SPRINGS', 'SHINGLE SPRINGS', 'COOL', 'WALNUT GROVE',
            'GARDEN VALLEY', 'SLOUGHHOUSE', 'WEST SACRAMENTO'], dtype=object)
```

```
[32]: # Other possible typo-error that can be checked are whitespace, fullstop, among
↳ others
new_data['city'] = new_data['city'].str.strip() # delete whitespace.
new_data['city'] = new_data['city'].str.replace('\.', '') # delete dot/full
↳ stop.
print(new_data.city.unique())
new_data.city.unique()
```

```
[32]: array(['SACRAMENTO', 'RANCHO CORDOVA', 'RIO LINDA', 'CITRUS HEIGHTS',
        'NORTH HIGHLANDS', 'ANTELOPE', 'ELK GROVE', 'ELVERTA', 'GALT',
        'CARMICHAEL', 'ORANGEVALE', 'FOLSOM', 'MATHER', 'POLLOCK PINES',
        'GOLD RIVER', 'EL DORADO HILLS', 'RANCHO MURIETA', 'WILTON',
        'GREENWOOD', 'FAIR OAKS', 'CAMERON PARK', 'LINCOLN', 'PLACERVILLE',
        'MEADOW VISTA', 'ROSEVILLE', 'ROCKLIN', 'AUBURN', 'LOOMIS',
        'EL DORADO', 'PENRYN', 'GRANITE BAY', 'FORESTHILL',
        'DIAMOND SPRINGS', 'SHINGLE SPRINGS', 'COOL', 'WALNUT GROVE',
        'GARDEN VALLEY', 'SLOUGHHOUSE', 'WEST SACRAMENTO'], dtype=object)
```

[4] The solution to the ‘street’ column:

```
[22]: # To check the count of unique values in the column
new_data.street.nunique()
```

[22]: 979

```
[33]: # There is actually less to be done here because the unique values almost equal
      ↳ the number of observations
      # Therefore, one way to clean the data is to remove blanks, dots, abbreviate some
      ↳ words, etc
new_data['street'] = new_data['street'].str.strip() # delete blankspaces
new_data['street'] = new_data['street'].str.replace('\\.', '') # delete dot/
      ↳ full stop.
print(new_data.street.nunique())
```

979

```
[24]: #changing the datatypes after the corrections have been effected
datatype= {'price': int, 'zip': int, 'longitude':float, 'latitude':float}
new_data = new_data.astype(datatype)
new_data['sale_date'] = pd.to_datetime(new_data.sale_date, format='%Y-%m-%d %H:
      ↳ %M:%S')
print(new_data.dtypes)
```

street	object
city	object
zip	int32
state	object
beds	float64
baths	float64
sq__ft	float64
type	object
sale_date	datetime64[ns]
price	int32
latitude	float64

```
longitude          float64
dtype: object
```

## 13 PHASE 3: HANDLING THE MISSING VALUES

```
[25]: new_data.isnull().values.any()
```

```
[25]: False
```

There are no missing values in the refined data. However, there are 'zero' valued cells which could also mean that the missing values have been replaced with zero. If the zero values actually represent missing values. Then, there are a number of ways to handle this:

- (i) Single-Value Imputation(SI) which involves replacing the missing cells with a single value. It could be the mean, highest occurring values, among others.
- (ii) Multiple/Multivariate Imputation(MI) which involves the use of different values to replace the missing cell based on the distribution of the data. There are several state of the art methods to do this.

My master thesis research was based on Classification with data irregularities (missing values and class imbalance). I implemented and compared different state of the art imputation algorithms such as Generative Adversarial Network (GAN) for building prediction. This could be a good alternative to handling the missing values. The link to my thesis can be found here <https://github.com/busarieedris/Classification-with-Data-Irregularities> (There may be some restrictions on some data due to privacy concerns. It is a collaborative research with a foremost research institute in Germany)

## 14 QUESTION THREE (3)

## 15 DATA SAVING

```
[26]: # Save the cleaned data with a better interactive name. This can be done with
      ↪ the '.to_csv' command
      # But the instruction says 'write a new csv with a similar name with the
      ↪ cleaned data'. That is the reason for changing the cleaned data
      # with a better name first.
      clean_realestate_fraugster_case=new_data.copy()
      clean_realestate_fraugster_case.to_csv('clean_realestate_fraugster_case.
      ↪ csv', index=False, sep=';')
```

```
[27]: clean_realestate_fraugster_case.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 983 entries, 0 to 984
Data columns (total 12 columns):
```

#	Column	Non-Null Count	Dtype
0	street	983 non-null	object
1	city	983 non-null	object
2	zip	983 non-null	int32
3	state	983 non-null	object
4	beds	983 non-null	float64
5	baths	983 non-null	float64
6	sq__ft	983 non-null	float64
7	type	983 non-null	object
8	sale_date	983 non-null	datetime64[ns]
9	price	983 non-null	int32
10	latitude	983 non-null	float64
11	longitude	983 non-null	float64

dtypes: datetime64[ns](1), float64(5), int32(2), object(4)  
memory usage: 92.2+ KB

## 16 QUESTION FOUR (4)

### 17 OUTPUT THE FOLLOWING ANSWERS:

(A) what is the distance (in meters) between the cheapest sale and the most recent sale?

SOLUTION / APPROACH

To do this:

STEP 1: You need the location (Longitude and Latitude) of the two points (The cheapest sale and the most recent sale).

```
[28]: # LET X BE THE CHEAPEST SALE (i.e The least value in the 'price' column)
lon_x=new_data.loc[new_data['price'].idxmin()]['longitude'] # The corresponding
↳ longitude for X
lat_x=new_data.loc[new_data['price'].idxmin()]['latitude'] # The corresponding
↳ latitude for X
```

```
[29]: # LET Y REPRESENT THE MOST RECENT SALE (i.e The most recent date in the
↳ 'sale_date' column)
lon_y=new_data.loc[new_data.sale_date.idxmax()]['longitude'] # The
↳ corresponding longitude for the most recent sale
lat_y=new_data.loc[new_data.sale_date.idxmax()]['latitude'] # The
↳ corresponding latitude for the most recent sale
```

STEP 2: Calculate the difference in distance between these two points

In order to get the distance between two coordinate points, there are quite some formulars for such calculations with varying degree of accuracy.

Some of the methods are:

- 1) Haversine formula: It is used to determine the distance between two points based on the law of Haversine.
- 2) Vincenty Formula: It is a distance calculation based on the fact that the earth is oblate spherical. It has an accuracy of almost 1mm

Step (i): Converting the trigonometrical values of the longitude and latitude into radian.

Step (ii): Find the difference in the coordinates.

Step (iii): Use one of the formulas above to calculate the distance between two points.

```
[34]: import math
from math import sin, cos, sqrt, atan2, radians
R = 6373.0 # Mean Radius of the Earth

# Step(i) Converting the trigonometrical values of the longitude and latitude
↳ into radian.
lat_x_rad = math.radians(lat_x)
lon_x_rad = math.radians(lon_x)
lat_y_rad = math.radians(lat_y)
lon_y_rad = math.radians(lon_y)

# Step(ii) Find the difference in the coordinates.
diff_lon = lon_y_rad - lon_x_rad
diff_lat = lat_y_rad - lat_x_rad

# Step(iii) For the purpose of this assignment, the Haversine formula would be
↳ used.
# Using Haversine formula to calculate the distance between two points.
a = math.sin(diff_lat / 2)**2 + math.cos(lat_x_rad) * math.cos(lat_y_rad) *
↳ math.sin(diff_lon / 2)**2
c = 2 * math.atan2(math.sqrt(a), math.sqrt(1 - a))
dist = R * c

print("The distance (in meters) between the cheapest sale and the most recent
↳ sale:", dist* 1000, 'metres')
```

The distance (in meters) between the cheapest sale and the most recent sale:  
6233.4089687788455 metres

.

- (B) what is the median street number, in multi-family houses, sold between 05/11/1933 and 03/12/1998, in Sacramento?

SOLUTION / APPROACH

To do this:

STEP 1: Filter out the rows with 'city= SACRAMENTO' and 'type= MULTI-FAMILY'

```
[35]: # Filter out the rows with 'city= SACRAMENTO' and 'type= MULTI-FAMILY' ,
data_add=new_data[(new_data['type']=='MULTI-FAMILY') &
↳(new_data['city']=='SACRAMENTO')]
data_add
```

```
[35]:
```

	street	city	zip	state	beds	baths	sq__ft	\
56	4520 BOMARK WAY	SACRAMENTO	95842	CA	4.0	2.0	1943.0	
68	7624 BOGEY CT	SACRAMENTO	95828	CA	4.0	4.0	2162.0	
108	2912 NORCADE CIR	SACRAMENTO	95826	CA	8.0	4.0	3612.0	
113	10158 CRAWFORD WAY	SACRAMENTO	95827	CA	4.0	4.0	2213.0	
353	2820 DEL PASO BLVD	SACRAMENTO	95815	CA	4.0	2.0	1404.0	
366	7342 GIGI PL	SACRAMENTO	95828	CA	4.0	4.0	1995.0	
527	398 LINDLEY DR	SACRAMENTO	95815	CA	4.0	2.0	1744.0	
648	8198 STEVENSON AVE	SACRAMENTO	95828	CA	6.0	4.0	2475.0	
716	1139 CLINTON RD	SACRAMENTO	95825	CA	4.0	2.0	1776.0	
923	7351 GIGI PL	SACRAMENTO	95828	CA	4.0	2.0	1859.0	

	type	sale_date	price	latitude	longitude
56	MULTI-FAMILY	1923-09-30 03:05:24	179580	38.665724	-121.358576
68	MULTI-FAMILY	2002-09-29 22:46:40	195000	38.480090	-121.415102
108	MULTI-FAMILY	2000-07-17 06:06:41	282400	38.559505	-121.364839
113	MULTI-FAMILY	1953-04-21 02:32:12	297000	38.570300	-121.315735
353	MULTI-FAMILY	1944-01-09 17:10:33	100000	38.617718	-121.440089
366	MULTI-FAMILY	2016-02-06 00:58:59	120000	38.490704	-121.410176
527	MULTI-FAMILY	1984-11-07 08:49:22	416767	38.622359	-121.457582
648	MULTI-FAMILY	1915-11-08 07:50:42	159900	38.465271	-121.404260
716	MULTI-FAMILY	1910-02-14 13:44:04	221250	38.585291	-121.406824
923	MULTI-FAMILY	1932-02-20 19:07:20	170000	38.490606	-121.410173

STEP 2: Filter the date that falls between '05/11/1933' and '03/12/1998' in step 1

```
[36]: # From the data_add gotten above, filter the date that falls between 05/11/1933,
↳and 03/12/1998
date_filter = (data_add['sale_date'] > '1933-11-05 00:00:00') &
↳(data_add['sale_date'] <= '1998-12-03 00:00:00') # Filter date 05/11/1933,
↳and 03/12/1998
data_ctd= data_add.loc[date_filter] # data with filtered city= SACRAMENTO,
↳type=MULTI-FAMILY and date=05/11/1933 and 03/12/1998.
data_ctd
```

```
[36]:
```

	street	city	zip	state	beds	baths	sq__ft	\
113	10158 CRAWFORD WAY	SACRAMENTO	95827	CA	4.0	4.0	2213.0	
353	2820 DEL PASO BLVD	SACRAMENTO	95815	CA	4.0	2.0	1404.0	
527	398 LINDLEY DR	SACRAMENTO	95815	CA	4.0	2.0	1744.0	

  

	type	sale_date	price	latitude	longitude
113	MULTI-FAMILY	1953-04-21 02:32:12	297000	38.570300	-121.315735

```
353 MULTI-FAMILY 1944-01-09 17:10:33 100000 38.617718 -121.440089
527 MULTI-FAMILY 1984-11-07 08:49:22 416767 38.622359 -121.457582
```

STEP 3: From the 'street' column, extract the characters before the first blankspace. This corresponds to the street numbers. Then, find the median of these numbers

```
[37]: # Extract street numbers from the street column (by splitting the content of
      ↪ the column by blank spaces and extracting the first value)
      # The result is passed to the median value method
      street_num = (data_ctd['street'].apply(lambda x: x.split()[0])).median()
      print('The median street number, in multi-family houses, sold between 05/11/
      ↪ 1933 and 03/12/1998 , in Sacramento is: ',street_num)
```

The median street number, in multi-family houses, sold between 05/11/1933 and 03/12/1998 , in Sacramento is: 2820.0

.

(C) What is the city name, and its 3 most common zip codes, that has the 2nd highest amount of beds sold?

SOLUTION / APPROACH

To do this:

STEP 1: Get the name of the city that has the 2nd highest amount of beds sold This is achieved by summing the number of beds per city.

The name of the city with the second highest number of sold beds is gotten

```
[38]: # Step 1: Get the name of the city that has the 2nd highest amount of beds sold
      # This is achieved by
      k=new_data.groupby('city')['beds'].sum()
      k.nlargest(2).iloc[[-1]]
```

```
[38]: city
      ELK GROVE      383.0
      Name: beds, dtype: float64
```

STEP 2: Filter out the ELK GROVE rows from the original data since we established that ELK GROVE is the city of interest.

```
[39]: # Filter out ELK GROVE rows from the original data since we established that ELK
      ↪ GROVE is the city of interest.
      data_elk=new_data[(new_data['city']!='ELK GROVE')]
      data_elk
```

```
[39]:
```

	street	city	zip	state	beds	baths	\
30	5201 LAGUNA OAKS DR UNIT 140	ELK GROVE	95758	CA	2.0	2.0	
34	5201 LAGUNA OAKS DR UNIT 162	ELK GROVE	95758	CA	2.0	2.0	
42	8718 ELK WAY	ELK GROVE	95624	CA	3.0	2.0	



50	9417 SARA ST	ELK GROVE	95624	CA	3.0	2.0
66	7005 TIANT % WAY	ELK GROVE	95758	CA	3.0	2.0
..	...	...	...	...	...	...
964	10085 ATKINS DR	ELK GROVE	95757	CA	3.0	2.0
965	9185 CERROLINDA CIR	ELK GROVE	95758	CA	3.0	2.0
975	5024 CHAMBERLIN CIR	ELK GROVE	95757	CA	3.0	2.0
979	1909 YARNELL WAY	ELK GROVE	95758	CA	3.0	2.0
983	8304 RED FOX WAY	ELK GROVE	95758	CA	4.0	2.0

	sq_ft	type	sale_date	price	latitude	longitude
30	1039.0	CONDO	1995-06-27 11:07:11	133000	38.423251	-121.444489
34	1039.0	CONDO	1995-02-14 02:55:54	141000	38.423251	-121.444489
42	1056.0	RESIDENTIAL	1964-12-06 07:32:56	156896	38.416530	-121.379653
50	1188.0	RESIDENTIAL	1913-11-26 14:12:17	170000	38.415518	-121.370527
66	1586.0	RESIDENTIAL	1911-11-05 06:37:27	194000	38.422811	-121.423285
..	...	...	...	...	...	...
964	1302.0	RESIDENTIAL	1965-08-27 15:55:54	219794	38.390893	-121.437821
965	1418.0	RESIDENTIAL	2003-11-28 09:52:55	220000	38.424497	-121.426595
975	1450.0	RESIDENTIAL	1919-02-03 11:26:02	228000	38.389756	-121.446246
979	1262.0	RESIDENTIAL	1979-04-17 04:20:28	230000	38.417382	-121.484325
983	1685.0	RESIDENTIAL	1933-04-10 20:13:38	235301	38.417000	-121.397424

[114 rows x 12 columns]

STEP 3 Find the three (3) most common zip codes of the ELK GROVE city

Do a groupby of zip with the GROVE city. This gives all the unique zip codes belonging to ELK GROVE

Then, count the number of occurrences (using the size()) of the unique ELK GROVE's zip codes and rename the resulting column as frequency (using the reset\_index(name='frequency')

Rearrange the table in descending order (sort\_values by ascending=[0,1]) so that the most frequent is the topmost in the column.

Use the .head(3) to select the first three rows as the three most frequent

```
[40]: data_elk.groupby(['city', 'zip']).size().reset_index(name='frequency').
      ↪ sort_values(['frequency', 'zip'], ascending=[0,1]).groupby('city').head(3)
```

```
[40]:      city  zip frequency
2  ELK GROVE  95758         44
1  ELK GROVE  95757         36
0  ELK GROVE  95624         34
```

```
[41]: # Print the solution
stg='''Therefore, the city name, and the 3 most common zip codes, that has the
      ↪ 2nd highest amount of beds sold: \n
city name: ELK GROVE \n
```

```
Zip codes: 95758,95757 and 95624'''  
print(stg)
```

Therefore, the city name, and the 3 most common zip codes, that has the 2nd highest amount of beds sold:

city name: ELK GROVE

Zip codes: 95758,95757 and 95624